


# MPI (Message Passing Interface)

## Sources

- [Introduction to the Message Passing Interface \(MPI\) using C](#)
-  [MPI Foundation Course: 6 Hours!](#)
- [Using MPI with C](#)
- <https://github.com/Amagnum/Parallel-Dot-Product-of-2-vectors-MPI/blob/main/main.cpp>
- <https://mpitutorial.com/tutorials/mpi-scatter-gather-and-allgather/>
- <https://github.com/mpitutorial/mpitutorial/blob/gh-pages/tutorials/mpi-scatter-gather-and-allgather/code/avg.c>

## Overview

- A C / C++ / Fortran library
- Serves to provide an efficient, multi-platform API to program in
- Follows the *Distributed Memory* paradigm
  - No shared variables
  - Processes communicate with each other by passing messages back and forth
- Header file
  - `#include <mpi.h>`
- Command to compile
  - `mpicc program.c -o program`
- Command to run (with runtime command line arguments)
  - `mpirun -np [num_processes] program`

## Handles

### Communicators

- Collection of process threads
- Each thread is assigned a unique ID, called it's *rank*
- In the beginning, all threads are grouped into one communicator
  - `MPI_COMM_WORLD`

### MPI Type Handles

- Similar to regular data types, used for the same purpose when passing messages using MPI
- Common handles include
  - `MPI_INT`

- MPI\_LONG
- MPI\_FLOAT
- MPI\_DOUBLE

## Typical Program Structure

```
int main(int argc, char **argv)
{
    // Variables in MPI aren't shared, so each process has its own copy of these
    int process_rank, process_count;

    // Spawn processes, command line arguments used to specify number of processes
    MPI_Init(&argc, &argv);

    // Get rank and process count
    MPI_Comm_rank(MPI_COMM_WORLD, &process_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &process_count);

    // Check if current process is root (master)
    if(process_rank == 0)
    {
        // Ask for user input
        // Split workload among threads, including itself
        // Perform its task and generate partial output
        // Collect partial outputs
        // Combine partial outputs to yield final output
    }

    // If not, it is a slave process
    else
    {
        // Perform its task and generate partial output
        // Send partial output to root
    }

    MPI_Finalize();
}
```

## Functions

### General Syntax

```
error_code = MPI_Xxxx(args...);
```

## Library Initialisation

Must be the first call to the MPI Library, spawns the processes

```
int MPI_Init(int *argc, char *argv);
```

## Library Finalisation

Must be the last call to the MPI Library, shuts down all threads except the root (rank 0)

```
int MPI_Finalize();
```

## Get Process Rank

```
MPI_Comm_rank(MPI_Comm comm, int *rank);
```

Eg.

```
int rank;
```

```
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
```

```
printf("Rank is %d\n", rank);
```

## Get Communicator Size

```
MPI_Comm_size(MPI_Comm comm, int *size);
```

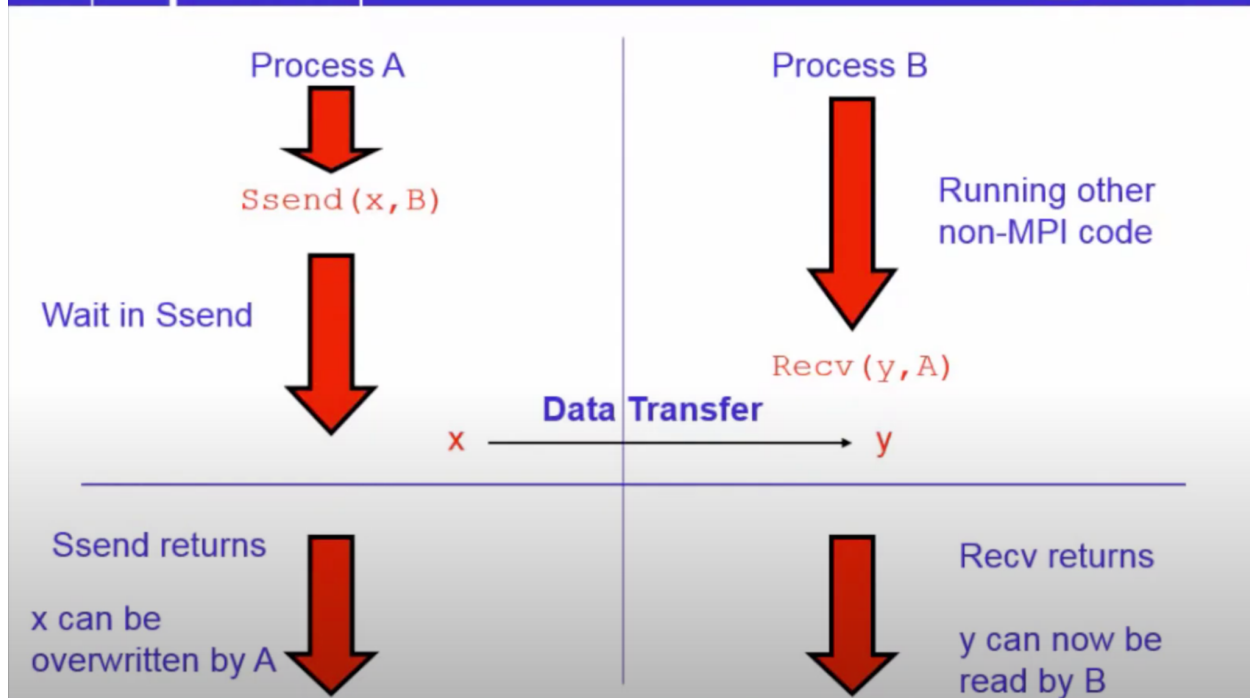
## Send and Recv

Send

Ssend

- (Synchronous send)
- Waits until the message is received before proceeding
- Prone to deadlocks, if both processes attempt to send at the same time

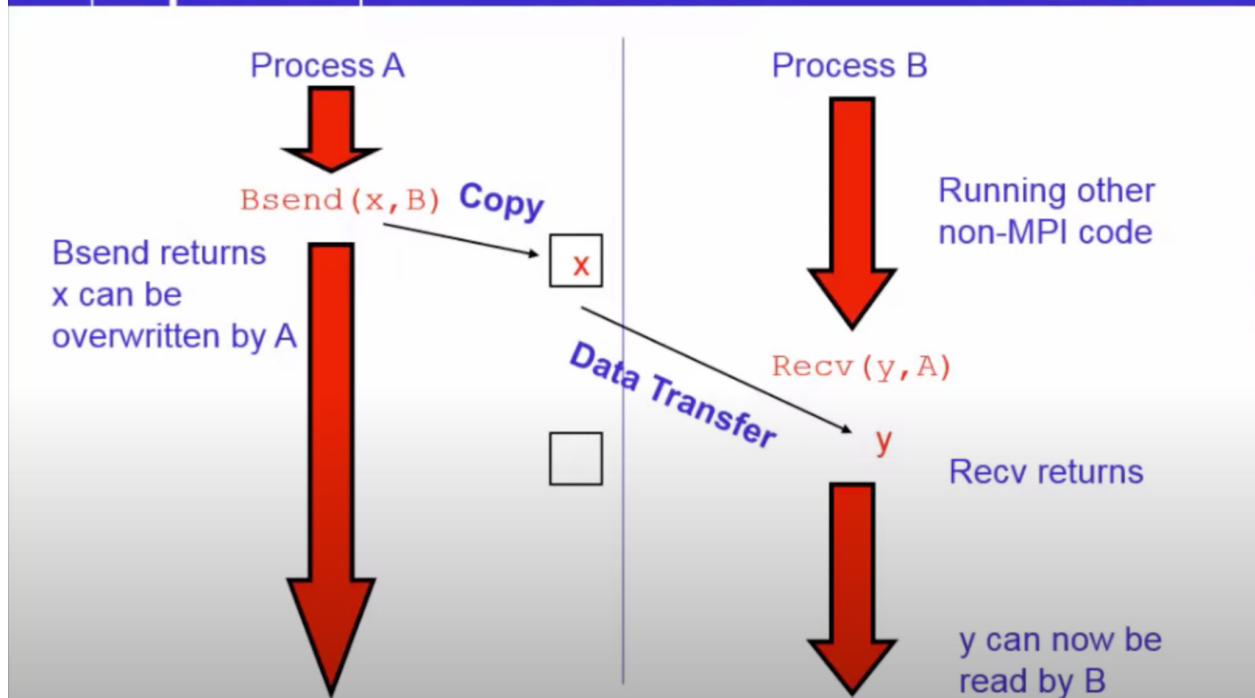
Eg.



#### Bsend

- (Buffered send)
- Asynchronous in nature
- Stores the message in a buffer, then proceeds with execution
- Sends the message when the receiving process runs `recv`
- Each process has a buffer space allocated to it for this purpose (size of 0 by default)
- Not prone to deadlocks
- Fails if buffer space is exhausted

Eg.



## Send

- Defaults to Bsend if buffer space is available
- If not, switches to Ssend
  - Vulnerable to deadlocks while in this state

## Syntax

```
int MPI_Send(...arguments);
```

## The arguments

- (void \*) *data\_to\_send*: address of a C variable that corresponds to *send\_type* below
- (int) *send\_count*: number of data elements to be sent
- (MPI type handle) *send\_type*: datatype of *data\_to\_send*
- (int) *destination\_ID*: rank of the receiving process
- (int) *tag*: message tag
- (MPI\_Comm) *comm*: communicator

## Recv

Only one function, acts in accordance with the type of send invoked

## Syntax

```
int MPI_Recv(...arguments);
```

The arguments

- (void \*) *received\_data*: address of a C variable that corresponds to *receive\_type* below
- (int) *receive\_count*: number of data elements to be received
- (MPI type handle) *receive\_type*: datatype of *received\_data*
- (int) *sender\_ID*: rank of the sending process
- (int) *tag*: message tag
- (MPI\_Comm) *comm*: communicator
- (MPI\_Status \*) *status*: metadata corresponding to *received\_data*

*receive\_count*, *sender\_ID* and *tag* can be used to select which message is to be received

### Example

A program to send a number from one process to another

```
#include<stdio.h>
#include<mpi.h>

int main(int argc, char **argv)
{
    int process_rank, process_count, num;
    MPI_Status status;

    MPI_Init(&argc, &argv);

    MPI_Comm_rank(MPI_COMM_WORLD, &process_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &process_count);

    if(process_rank == 0)
    {
        printf("Enter the value of num\n");
        scanf("%d", &num);

        MPI_Send(&num, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
    }

    else
    {
        MPI_Recv(&num, 1, MPI_INT, 0, 1, MPI_COMM_WORLD, &status);

        printf("Received %d by process %d", num, process_rank);
    }

    MPI_Finalize();
}
```

```
}
```

## Barrier

Used for synchronisation, blocks progression until all processes have reached the barrier

### Syntax

```
MPI_Barrier();
```

## Broadcast, Scatter and Gather

These are collective operations that involve all processes

### Broadcast

- One call, all must execute.
- No separate receive function.
- Other than the sender, all processes wait to receive and synchronise

### Syntax

```
int MPI_Bcast(...arguments);
```

#### The arguments

- (void \*) *data\_to\_send*: address of a C variable that corresponds to *send\_type* below
- (int) *send\_count*: number of data elements to be sent
- (MPI type handle) *send\_type*: datatype of *data\_to\_send*
- (int) *sender\_ID*: rank of the sending (broadcasting) process
- (MPI\_Comm) *comm*: communicator

### Scatter

- Split a variable into subsets of values
- Distribute them among processes in the group defined by the communicator
- Works similarly to broadcast

### Syntax

```
int MPI_Scatter(...arguments); // for a fixed subset size
```

```
int MPI_Scatterv(...arguments); // for a variable subset size (data not perfectly divisible by the number of processes)
```

#### The arguments

- (void \*) *data\_to\_scatter*: address of a C variable that corresponds to *send\_type* below
- (int) *send\_count*: number of data elements to send to each process
- (MPI type handle) *send\_type*: datatype of *data\_to\_scatter*
- (void \*) *data\_to\_receive*: subset of *data\_to\_scatter*
- (int) *receive\_count*: size of the subset in *data\_to\_receive*

- (MPI type handle) *receive\_type*: datatype of *data\_to\_receive*
- (int) *sender\_id*: rank of sending process
- (MPI\_Comm) *comm*: communicator

## Gather

- Companion function to *Scatter*
- Gathers values from all sent variables into a single variable

## Syntax

int MPI\_Gather(...arguments); // for a fixed subset size

int MPI\_Gatherv(...arguments); // for a variable subset size (data not perfectly divisible by the number of processes)

## The arguments

- (void \*) *data\_to\_gather*: address of the C variable to be sent by all processes that corresponds to *send\_type* below
- (int) *send\_count*: number of data elements to be sent by each process
- (MPI type handle) *send\_type*: datatype of *data\_to\_scatter*
- (void \*) *gathering\_variable*: variable that holds all values of *data\_to\_gather*
- (int) *receive\_count*: size of the variable to be received from each process
- (MPI type handle) *receive\_type*: datatype of *data\_to\_receive*
- (int) *receiver\_id*: rank of receiving process
- (MPI\_Comm) *comm*: communicator

## Reduce / Allreduce

- Reduce
  - Used to merge partial results by employing an operator
  - Places final result in the root process (rank = receiver\_id)
  - Common operators include
    - MPI\_SUM
    - MPI\_PROD
    - MPI\_MAX
    - MPI\_MIN
- All\_Reduce
  - Combines partial results into a variable held by all participating processes
  - Omit the last but one argument (receiver\_id)

## Syntax

int MPI\_Reduce(...arguments); // for a variable subset size

## The arguments

- (void \*) *data\_to\_reduce*: address of a C variable that corresponds to *send\_type* below, that contains a partial result to be merged
- (void \*) *reduced\_result*: address of a C variable that will contain the final result of the merge



- (int) *send\_count*: number of data elements to send
- (MPI type handle) *send\_type*: datatype of *data\_to\_reduce*
- (MPI operation handle) *operation\_type*: type of operation to be performed
- (int) *receiver\_id*: rank of receiving (destination of final result) process, typically the root
- (MPI\_Comm) *comm*: communicator

### Example

A program to compute the dot product of two vectors

```
#include<stdio.h>
#include<mpi.h>
#define MAX 10000

int main(int argv, char **argc)
{
    int process_rank, process_count, n, s_size, sum, dot_p;
    int a[MAX], b[MAX], a_s[MAX], b_s[MAX];

    MPI_Init(&argv, &argc);

    MPI_Comm_rank(MPI_COMM_WORLD, &process_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &process_count);

    // Root process asks user for input vectors
    if(process_rank == 0)
    {
        printf("Enter the length of the vectors\n");
        scanf("%d", &n);

        printf("Enter the elements of the first vector\n");
        for(int i = 0; i < n; i++)
        {
            scanf("%d", &a[i]);
        }

        printf("Enter the elements of the second vector\n");
        for(int i = 0; i < n; i++)
        {
            scanf("%d", &b[i]);
        }

        s_size = n / process_count;
    }
}
```

```

// Root process broadcasts length of vectors
MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);

// Root process broadcasts size of each subset
MPI_Bcast(&s_size, 1, MPI_INT, 0, MPI_COMM_WORLD);

// Root process scatters each vector among all processes
MPI_Scatter(a, s_size, MPI_INT, a_s, s_size, MPI_INT, 0, MPI_COMM_WORLD);
MPI_Scatter(b, s_size, MPI_INT, b_s, s_size, MPI_INT, 0, MPI_COMM_WORLD);

// Each process computes its partial sum
for(int i = 0; i < s_size; i++)
{
    sum += a_s[i] * b_s[i];
}

// Combining the partial sums
MPI_Reduce(&sum, &dot_p, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);

if(process_rank == 0)
{
    printf("The dot product is %d\n", dot_p);
}

MPI_Finalize();
}

```

## Lab Programs

### Trapezoidal Rule

Write an MPI program to compute the area under the curve using trapezoidal rule using MPI\_Reduce and MPI\_Allreduce.

```

#include<stdio.h>
#include<mpi.h>

// Function to be integrated over
double f(double x)
{
    return (1 / (1 + (x * x)));
}

double trapezoidal(double a_p, double b_p, double h)

```

```

{
    // For each interval, i is the lower, and j is the upper bound
    double i, j;
    double y_i, y_j, area, sum_p = 0.0;

    for(i = a_p; i < b_p; i += h)
    {
        j = i + h;

        y_i = f(i);
        y_j = f(j);

        // Area of trapezium defined by this interval
        area = 0.5 * h * (y_i + y_j);

        // Calculate partial sum (sub-integral)
        sum_p += area;
    }

    return sum_p;
}

int main(int argc, char **argv)
{
    double a, b, n, a_p, b_p, n_p, sum, sum_p, h;

    int process_rank, num_processes;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &process_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &num_processes);

    // Only master process receives input from user
    if(process_rank == 0)
    {
        printf("Enter a, b and n\n");
        scanf("%lf", &a);
        scanf("%lf", &b);
        scanf("%lf", &n);

        // Number of intervals assigned to each process
        n_p = n / num_processes;

        // Length of each interval

```

```

    h = (b - a) / n;
}

// Master process broadcasts user input to slave processes
MPI_Bcast(&a, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);
MPI_Bcast(&b, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);
MPI_Bcast(&n_p, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);
MPI_Bcast(&h, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);

// a = 0, b = 20, n = 10, h = 2, num_processes = 2, n_p = 5

a_p = a + process_rank * h * n_p;
b_p = a_p + h * n_p;

sum_p = trapezoidal(a_p, b_p, h);

// Combine partial sums
MPI_Reduce(&sum_p, &sum, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
MPI_Allreduce(&sum_p, &sum, 1, MPI_DOUBLE, MPI_SUM, MPI_COMM_WORLD);

printf("The area under the curve is %lf (Process %d)\n", sum, process_rank);

MPI_Finalize();
}

```

## Sum of Vectors

Write an MPI program to read 2 vectors and print the sum vector using MPI\_Scatter and MPI\_Gather.

$$\begin{aligned}
 \mathbf{x} + \mathbf{y} &= (x_0, x_1, \dots, x_{n-1}) + (y_0, y_1, \dots, y_{n-1}) \\
 &= (x_0 + y_0, x_1 + y_1, \dots, x_{n-1} + y_{n-1}) \\
 &= (z_0, z_1, \dots, z_{n-1})
 \end{aligned}$$

```

#include<stdio.h>
#include<mpi.h>

```

```

int main(int argc, char **argv)
{
    int process_rank, num_processes;
    int i, n, n_p;
    int a[100], b[100], sum[100], a_p[100], b_p[100], sum_p[100];

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &process_rank);

```

```

MPI_Comm_size(MPI_COMM_WORLD, &num_processes);

if(process_rank == 0)
{
    printf("Enter the length of the vectors\n");
    scanf("%d", &n);

    printf("Enter the elements of vector 1\n");
    for(i = 0; i < n; i++)
    {
        scanf("%d", &a[i]);
    }

    printf("Enter the elements of vector 2\n");
    for(i = 0; i < n; i++)
    {
        scanf("%d", &b[i]);
    }

    n_p = n / num_processes;
}

MPI_Bcast(&n_p, 1, MPI_INT, 0, MPI_COMM_WORLD);
MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);

MPI_Scatter(a, n_p, MPI_INT, a_p, n_p, MPI_INT, 0, MPI_COMM_WORLD);
MPI_Scatter(b, n_p, MPI_INT, b_p, n_p, MPI_INT, 0, MPI_COMM_WORLD);

for(i = 0; i < n_p; i++)
{
    sum_p[i] = a_p[i] + b_p[i];
}

MPI_Gather(sum_p, n_p, MPI_INT, sum, n_p, MPI_INT, 0, MPI_COMM_WORLD);

if(process_rank == 0)
{
    printf("Resultant vector :-\n");

    for(i = 0; i < n; i++)
    {
        printf("%d ", sum[i]);
    }
}

```

```

    printf("\n");
}

MPI_Finalize();
return 0;
}

```

## Get Input

Write a function `get_input(int rank, int comm_size, double *a, double *b, int *n)` to read 3 values viz. `a` (double), `b` (double) and `n` (int) on process 0 and send it to other processes. Rewrite the same function using `MPI_Bcast()` method.

```

#include<stdio.h>
#include<mpi.h>

```

```

void get_input(double *a, double *b, int *n, int process_rank, int num_processes, int b_flag)
{
    MPI_Status status;
    double a_val = 0.0, b_val = 0.0;
    int n_val = 0;

    if(process_rank == 0)
    {
        printf("Enter a, b and n\n");
        scanf("%lf", &a_val);
        scanf("%lf", &b_val);
        scanf("%d", &n_val);

        a = &a_val;
        b = &b_val;
        n = &n_val;
    }

    if(b_flag == 0)
    {
        if(process_rank == 0)
        {
            for(int i = 1; i < num_processes; i++)
            {
                MPI_Send(a, 1, MPI_DOUBLE, i, 0, MPI_COMM_WORLD);
                MPI_Send(b, 1, MPI_DOUBLE, i, 1, MPI_COMM_WORLD);
                MPI_Send(n, 1, MPI_INT, i, 2, MPI_COMM_WORLD);
            }
        }
    }
}

```

```

    }
    else
    {
        MPI_Recv(a, 1, MPI_DOUBLE, 0, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
        MPI_Recv(b, 1, MPI_DOUBLE, 0, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
        MPI_Recv(n, 1, MPI_INT, 0, MPI_ANY_TAG, MPI_COMM_WORLD, &status);

        printf("a = %lf (Process - %d)\n", *a, process_rank);
        printf("b = %lf (Process - %d)\n", *b, process_rank);
        printf("n = %d (Process - %d)\n", *n, process_rank);
    }
}

else
{
    MPI_Bcast(a, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);
    MPI_Bcast(b, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);
    MPI_Bcast(n, 1, MPI_INT, 0, MPI_COMM_WORLD);

    if(process_rank != 0)
    {
        printf("a = %lf (Process - %d)\n", *a, process_rank);
        printf("b = %lf (Process - %d)\n", *b, process_rank);
        printf("n = %d (Process - %d)\n", *n, process_rank);
    }
}

}

int main(int argc, char **argv)
{
    double a, b;
    int n, b_flag;

    int process_rank, num_processes;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &process_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &num_processes);

    if(process_rank == 0)
    {
        printf("Enter 1 to broadcast the values, 0 to not\n");
    }
}

```

```
    scanf("%d", &b_flag);  
}  
  
// Broadcasting b_flag alone  
MPI_Bcast(&b_flag, 1, MPI_INT, 0, MPI_COMM_WORLD);  
  
get_input(&a, &b, &n, process_rank, num_processes, b_flag);  
  
MPI_Finalize();  
}
```